

# ELASTIC/ ELK STACK

- **E – Elasticsearch**
- **L- Logstash**
- **K- Kibana**

## ElasticSearch:

- It is open-source **Search & Analytics Engine** and can also serve as a **NOSQL Database** which will store data in the form of **Json** and uses **RESTFUL API** to store and retrieve data.
- It works based on Apache Lucene which is also known as “**Heart of Elasticsearch**”

## Logstash:

It is used to **read, write, filter and modify data** from various sources and store it in Elasticsearch.

## Kibana:

- It is a web-interface which is used to **Discover, Analyze, Monitor and visualize** the data from Elasticsearch.
- It also used to apply **Machine Learning Algorithms** on the data from Elasticsearch to get insights of **data Anomaly** and **future trends**.

## Beats:

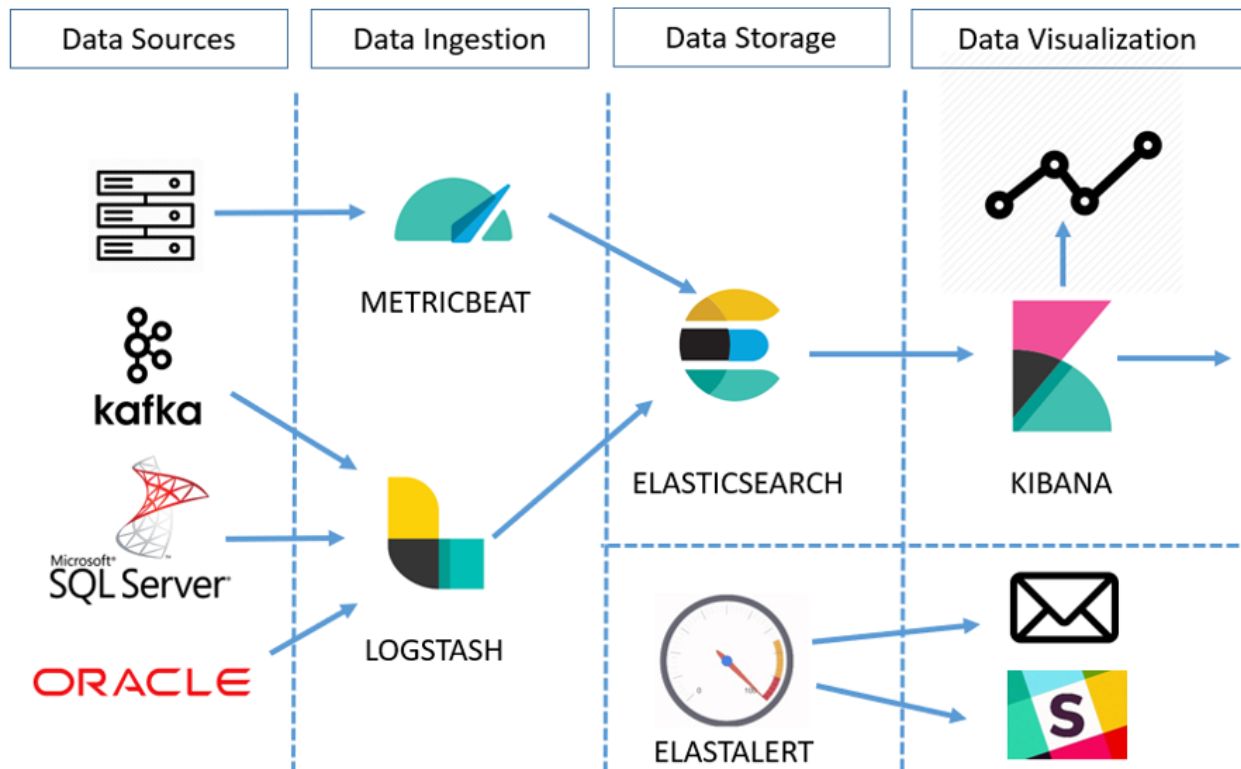
They are Light-weight data shippers which is used to ship data from data source to Elasticsearch.

## **Types:**

File Beats	Log files
Metric Beats	Metrics(CPU,Memory)
Packet Beats	Network Data
Win log Beats	Windows Event Logs
Audit Beats	Audit (OS files)
Heart Beats	Uptime Monitor
Function Beats	Serverless Shipper

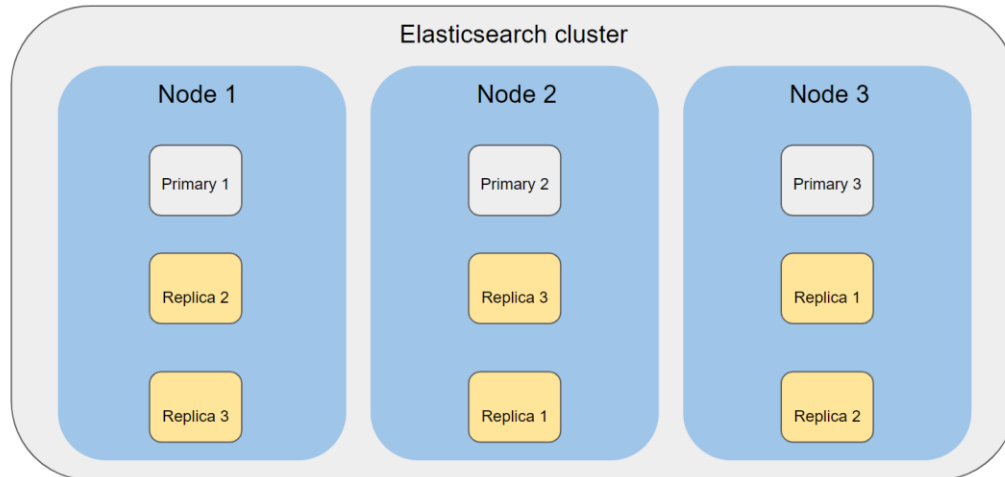
## Architecture Of Elastic Stack:

### Monitoring Architecture Overview



## Terminology in Elastic Stack:

Elastic terms	RDBMS terms	Refers to
Fields	Columns	Key-value Pair (JSON Objects)
Documents	Rows	Collection of Fields
Index	Table	Collection of Documents
Cluster	Database	Collection of Index
Shards	-	Horizontal partitioning of Index
Replica	-	Copy of Shards



Primary Shards and its Replicas in Elasticsearch cluster

### Types of Nodes in Cluster:

Master Node	<ul style="list-style-type: none"> <li>• Responsible for creation or Deletion of Index.</li> <li>• Tracks the other nodes.</li> <li>• Determines the location of shards.</li> </ul>
Data Node	Responsible for performing CRUD, Search and Aggregation functions.
Ingest Node	Responsible for processing a document before indexing them.
Co-ordinating Node	<ul style="list-style-type: none"> <li>• Performs Routing</li> <li>• Aids for Search Reduction Phase</li> <li>• Responsible for Distributing the works via BULK Indexing.</li> </ul>

### Installation and Set Up Procedure:

#### Server:

- Install Oracle VM.
- Install Ubuntu. (Server)
- Download and install:
  - `wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo gpg --dearmor -o /usr/share/keyrings/elasticsearch-keyring.gpg`
  - `sudo apt-get install apt-transport-https`
  - `echo "deb [signed-by=/usr/share/keyrings/elasticsearch-keyring.gpg] https://artifacts.elastic.co/packages/8.x/apt stable main" | sudo tee /etc/apt/sources.list.d/elastic-8.x.list`
  - `sudo apt-get update && sudo apt-get install elasticsearch`

- sudo nano /etc/elasticsearch/elasticsearch.yml  
After this enable the following fields in the .yaml file:
  - node.name: node-1
  - network.host: 0.0.0.0
  - discovery.seed.hosts: ["127.0.0.1"]
  - xpack.security.enabled: false
  - cluster.initial\_master\_nodes: ["node-1"]
- sudo chmod 755 -R /var/log/elasticsearch/
- Configure Elasticsearch to start automatically when the system boots up
  - sudo /bin/systemctl daemon-reload
  - sudo /bin/systemctl enable elasticsearch.service
- Elasticsearch can be started as follows
  - sudo /bin/systemctl start elasticsearch.service
  - sudo /bin/systemctl status elasticsearch.service
- Install curl
  - sudo apt-get install curl
- Go to Settings ->Networks ->Port Forwarding  
and add the following Network Configuration:
  - **Elasticsearch-127.0.0.1-9200**
  - **Kibana 127.0.0.1 5601**
  - **SSL 127.0.0.1 22**
- Open Terminal and Type the following Commands:
  - Sudo apt-get install openssh-server
  - Sudo systemctl enable ssh
  - Sudo adduser username
  - Sudo usermod -aG sudo username

### Client:

- Download Putty. (Client)
- Configure the ElasticSearch in 127.0.0.1 in port 9200 and load it and open it.
- Login as username and enter the password.
- Then load the dataset into Elasticsearch using the Command:
  - Dataset: <http://media.sundog-soft.com/es8/movies.json>
- Create the index and post the data using the command:
  - curl -X PUT "localhost:9200/movies?pretty"
  - curl -XPOST "localhost:9200/movies/\_bulk?pretty" --data-binary @movies.json

### CRUD OPERATIONS IN INDEX:

- Command to know the mappings of the index movies:
  - curl -XGET "127.0.0.1:9200/movies/\_mappings?pretty" (Note: To reduce the mapping area of any index we can define the datatype as "Flattened")
- Command to add document to the index movies:

- curl -XPUT 127.0.0.1:9200/movies/\_doc/
- Command to delete document from the index movies:
  - curl -XDELETE 127.0.0.1:9200/movies/\_doc/1234567 //1234567-id
- Command to Update a document to the index movies:
  - Curl -XPUT 127.0.0.1:9200/movies/\_doc/1234567
- Command to do OPTIMISTIC CONCURRENCY CONTROL to update a document is:
  - Curl -XPUT 127.0.0.1:9200/movies/\_doc/1234567?if\_seq\_no=7&if\_primary\_term=1

## **REALTIONSHIP BETWEEN DOCUMENTS IN A INDEX:**

### **Command to establish a parent child relationship between franchise and film:**

```
Curl -XPUT 127.0.0.1:9200/series -d ' {
  "mappings" : {
    "properties" : {
      "film_to_franchise" : {
        "type" : "join" ,
        "relations": { "franchise" , "flim"}           //franchise -parent
      }                                               //flim-child
    }
  }
}'
```

### **Command to find the child who has "franchise" as parent:**

```
Curl -XGET 127.0.0.1:9200/series/_search?pretty -d' {
  "query" : {
    "has_parent" : {
      "parent_type": "franchise",
      "query" : { "match" : { "title" : "Star Wars"} }
    }
  }
}'
```

```

student@harsini-VirtualBox:~$ curl -XGET 127.0.0.1:9200/series/_search?pretty -d' {
  "query" : {
    "has_parent" : {
      "parent_type": "franchise",
      "query" : { "match" : { "title" : "Star Wars" } }
    }
  }
},
{
  "took" : 3,
  "timed_out" : false,
  "_shards" : {
    "total" : 1,
    "successful" : 1,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : {
      "value" : 7,
      "relation" : "eq"
    },
    "max_score" : 1.0,
    "hits" : [
      {
        "_index" : "series",
        "_id" : "260",
        "_score" : 1.0,
        "_routing" : "1",
        "_source" : {
          "id" : "260",
          "film_to_franchise" : {
            "name" : "film",
            "parent" : "1"
          },
          "title" : "Star Wars: Episode IV - A New Hope",
          "year" : "1977",
          "genre" : [
            "Action",
            "Adventure",
            "Sci-Fi"
          ]
        }
      }
    ]
  },
},

```

### **Command to find the parent who has “The Force Awakens” as child:**

Curl -XGET 127.0.0.1:9200/series/\_search?pretty -d' {

```

  "query" : {
    "has_child" : {
      "type": "film",
      "query" : { "match" : { "title" : "The Force Awakens" } }
    }
  }
},

```

```

student@harsini-VirtualBox:~$ curl -XGET 127.0.0.1:9200/series/_search?pretty -d' {
  "query" : {
    "has_child" : {
      "type": "film",
      "query" : { "match" : { "title" : "The Force Awakens" } }
    }
  }
}'
{
  "took" : 104,
  "timed_out" : false,
  "_shards" : {
    "total" : 1,
    "successful" : 1,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : {
      "value" : 1,
      "relation" : "eq"
    },
    "max_score" : 1.0,
    "hits" : [
      {
        "_index" : "series",
        "_id" : "1",
        "_score" : 1.0,
        "_routing" : "1",
        "_source" : {
          "id" : "1",
          "film_to_franchise" : {
            "name" : "franchise"
          },
          "title" : "Star Wars"
        }
      }
    ]
  }
}

```

## **SEARCH IN ELASTICSEARCH:**

### **Query Line Search**

The query is given directly as a parameter

Example: Query to get the details of the movie which is released after the year 2010 and has the word “trek” in the title.

- Curl -XGET ‘http://localhost:9200/movies/\_search?q=+year:>2010+title:trek’

### **Request Body Search**

The query is given as a request body

Example: Query to get the details of the movie which is released after the year 2010 and has the word “trek” in the title.

- Curl -XGET ‘http://localhost:9200/movies/\_search?pretty’ -d ‘{
 “query”: {
 “bool”: {
 “must”: { “term” : { “title” : “trek”}},
 “filter”: { “range” : { “year” : { “gte” : 2010}}}

```
}  
}  
}'
```

```
"took" : 269,  
"timed_out" : false,  
"_shards" : {  
  "total" : 1,  
  "successful" : 1,  
  "skipped" : 0,  
  "failed" : 0  
},  
"hits" : {  
  "total" : {  
    "value" : 2,  
    "relation" : "eq"  
  },  
  "max_score" : 6.501652,  
  "hits" : [  
    {  
      "_index" : "movies",  
      "_id" : "135569",  
      "_score" : 6.501652,  
      "_source" : {  
        "id" : "135569",  
        "title" : "Star Trek Beyond",  
        "year" : 2016,  
        "genre" : [  
          "Action",  
          "Adventure",  
          "Sci-Fi"  
        ]  
      }  
    },  
    {  
      "_index" : "movies",  
      "_id" : "102445",  
      "_score" : 5.7094297,  
      "_source" : {  
        "id" : "102445",  
        "title" : "Star Trek Into Darkness",  
        "year" : 2013,  
        "genre" : [  
          "Action",  
          "Adventure",  
          "Sci-Fi",  
          "IMAX"  
        ]  
      }  
    }  
  ]  
}
```

## Difference between match and match\_phrase

### ☐ Term Matching:

- **match:** Breaks down the input text into individual terms and matches any of them.
- **match\_phrase:** Searches for the exact sequence of terms as a phrase.

### ☐ Order and Proximity:

- **match:** Ignores the order of terms and proximity.
- **match\_phrase:** Considers the order and ensures the terms appear close to each other as specified.



## Match Phrase Search

Search for exact phrases given in the query

```
➤ curl -XGET 'http://localhost:9200/movies/_search?pretty' -d '{
  "query": {
    "match_phrase": {
      "title": {
        "query": "Star Wars",
        "slop": 2
      }
    }
  }
}'
```

```
"timed_out" : false,
"_shards" : {
  "total" : 1,
  "successful" : 1,
  "skipped" : 0,
  "failed" : 0
},
"hits" : {
  "total" : {
    "value" : 13,
    "relation" : "eq"
  },
  "max_score" : 9.028755,
  "hits" : [
    {
      "_index" : "movies",
      "_id" : "61160",
      "_score" : 9.028755,
      "_source" : {
        "id" : "61160",
        "title" : "Star Wars: The Clone Wars",
        "year" : 2008,
        "genre" : [
          "Action",
          "Adventure",
          "Animation",
          "Sci-Fi"
        ]
      }
    },
    {
      "_index" : "movies",
      "_id" : "135216",
      "_score" : 9.028755,
      "_source" : {
        "id" : "135216",
        "title" : "The Star Wars Holiday Special",
        "year" : 1978,
        "genre" : [
          "Adventure",
          "Children",
          "Comedy",
          "Sci-Fi"
        ]
      }
    }
  ]
}
```

## Pagination

While searching for the query we can do pagination by defining from and size keywords.

- From – specifies the starting point
- Size – specifies the number of results to be retrieved

```
➤ curl -XGET 'http://localhost:9200/movies/_search?pretty' -d '{
    "query": {
        "match": {
            "title": "Star"
        }
    },
    "from": 0,
    "size": 10
}'
```

```
"hits" : [
  {
    "_index" : "movies",
    "_id" : "800",
    "_score" : 6.377307,
    "_source" : {
      "id" : "800",
      "title" : "Lone Star",
      "year" : 1996,
      "genre" : [
        "Drama",
        "Mystery",
        "Western"
      ]
    }
  },
  {
    "_index" : "movies",
    "_id" : "1613",
    "_score" : 6.377307,
    "_source" : {
      "id" : "1613",
      "title" : "Star Maps",
      "year" : 1997,
      "genre" : [
        "Drama"
      ]
    }
  }
],
}
```

## Sorting

It is used to sort the result which is fetched using search query

```
➤ Curl -XGET 127.0.0.1:9200/movies/_search?sort=year&pretty'
```

In order to sort based on text value we need to define them as keyword in the raw data format

```
➤ Curl -XPUT 127.0.0.1:9200/movies/ -d '{
  "mappings": {
    "properties": {
      "title": {
        "type": "text",
        "fields": { "raw": { "type": "keyword" } }
      }
    }
  }
}'
```

### Difference between Text and Keyword fields

#### Text Fields

- **Purpose:** Text fields are used for full-text search. They are analyzed, meaning the text is processed and broken down into individual terms (tokens) using an analyzer.

#### Keyword Fields

- **Purpose:** Keyword fields are used for exact matching, sorting, and aggregations. They are not analyzed, meaning the text is indexed as a single token.

### Why Have a Keyword Field for a Text Field?

1. **Exact Matching:** When you need to perform exact match queries on a field, such as finding all documents where the title is exactly "Star Wars." Analyzing the text would break it down into individual terms, making it unsuitable for exact matches.
2. **Sorting:** Sorting requires the exact values of the field. Analyzed text fields cannot be sorted properly because they are broken down into multiple terms.
3. **Aggregations:** Aggregations, like counting unique values, require the exact terms. Text fields, which are analyzed, cannot be used for accurate aggregations.

### Fuzzy Queries

Fuzzy queries are designed to handle search terms that may contain misspellings or typographical errors. They can identify similar terms within a certain edit distance, allowing for more flexible searches. Fuzzy queries support:

- Substitution Eg: Apple -> Appl
  - Insertion Eg: Apple -> Applea
  - Deletion Eg: Apple->Aple
- `curl -XGET 'http://localhost:9200/movies/_search?pretty' -d '{`
- ```

"query": {
  "fuzzy": {
    "title": {"value": "Ster", "fuzziness": 1}
  }
}
}
```

```
{
  "took" : 179,
  "timed_out" : false,
  "_shards" : {
    "total" : 1,
    "successful" : 1,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : {
      "value" : 38,
      "relation" : "eq"
    },
    "max_score" : 4.7829804,
    "hits" : [
      {
        "_index" : "movies",
        "_id" : "800",
        "_score" : 4.7829804,
        "_source" : {
          "id" : "800",
          "title" : "Lone Star",
          "year" : 1996,
          "genre" : [
            "Drama",
            "Mystery",
            "Western"
          ]
        }
      },
      {
        "_index" : "movies",
        "_id" : "1613",
        "_score" : 4.7829804,
        "_source" : {
          "id" : "1613",
          "title" : "Star Maps",
          "year" : 1997,
          "genre" : [
            "Drama"
          ]
        }
      }
    ]
  }
}
```

## **Partial Matching**

It is used to perform partial matching for search

### **Prefix query**

- `curl -XGET 'http://localhost:9200/movies/_search?pretty' -d '{`

```

        "query": {
            "prefix": {
                "year": "201"
            }
        }
    },
}
```

### **Wildcard query**

- `curl -XGET 'http://localhost:9200/movies/_search?pretty' -d '{`

```

    "query": {
        "wildcard": {
            "year": "19*"
        }
    }
}
```

### **Search-as-you-type**

As the name suggests, it will perform search as you type

- `curl -XGET 'http://localhost:9200/movies/_search?pretty' -d '{`

```

    "query": {
        "match_phrase_prefix": {
            "title": {
                "query": "Star Wars",
                "slop": 10
            }
        }
    }
}
```

```

"hits" : [
  {
    "_index" : "movies",
    "_id" : "61160",
    "_score" : 15.817083,
    "_source" : {
      "id" : "61160",
      "title" : "Star Wars: The Clone Wars",
      "year" : 2008,
      "genre" : [
        "Action",
        "Adventure",
        "Animation",
        "Sci-Fi"
      ]
    }
  },
  {
    "_index" : "movies",
    "_id" : "135216",
    "_score" : 15.817083,
    "_source" : {
      "id" : "135216",
      "title" : "The Star Wars Holiday Special",
      "year" : 1978,
      "genre" : [
        "Adventure",
        "Children",
        "Comedy",
        "Sci-Fi"
      ]
    }
  }
],

```

## **EXCEPTION HANDLING FOR SEARCH:**

### **To ignore the exception which is throw by datatype:**

Example: When we try to perform search operation using the keyword datatype:

```

Curl -location -XPUT 127.0.0.1:9200/microservice-logs/_settings \
--data-raw ' {
    "index.mapping.ignore_malformed" : true
}

```

### **To ignore the exception which is throw because we are exceeding the default limit(1000):**

```

Curl -location -XPUT 127.0.0.1:9200/big_Objects/_settings \
--data-raw ' {
    "index.mapping.total.fields.limit" : 1005
}

```

## **Importing Data from different sources to Elasticsearch**

- Java- Elastic.co
- Python – Elasticsearch.package
- Ruby-Elasticsearch.ruby
- Perl-Elasticsearch.pm

## **Importing Data using Python:**

Create and Run IndexRatings.py file

### **IndexRatings.py file:**

```
import csv
from collections import deque
import elasticsearch
from elasticsearch import helpers

def readMovies():
    csvfile = open('ml-latest-small/movies.csv', 'r', encoding="utf8")

    reader = csv.DictReader( csvfile )

    titleLookup = {}

    for movie in reader:
        titleLookup[movie['movieId']] = movie['title']

    return titleLookup

def readRatings():
    csvfile = open('ml-latest-small/ratings.csv', 'r', encoding="utf8")

    titleLookup = readMovies()

    reader = csv.DictReader( csvfile )
    for line in reader:
        rating = {}
        rating['user_id'] = int(line['userId'])
        rating['movie_id'] = int(line['movieId'])
        rating['title'] = titleLookup[line['movieId']]
        rating['rating'] = float(line['rating'])
        rating['timestamp'] = int(line['timestamp'])
        yield rating

es = elasticsearch.Elasticsearch(["http://127.0.0.1:9200"])

#es.indices.delete(index="ratings",ignore=404)
deque(helpers.parallel_bulk(es,readRatings(),index="ratings", request_timeout=300), maxlen=0)
es.indices.refresh()
```

### **Run the command**

```
curl -XGET 127.0.0.1:9200/ratings/_search?pretty
```

## **Importing Data from MySql:**

### **Install Mysql Connector:**

- `sudo apt-get install mysql-server`
- `wget http://files.grouplens.org/datasets/movielens/ml-100k.zip`
- `unzip ml-100k.zip`
- `sudo mysql --local-infile=1 -u root -p`
- `CREATE DATABASE movielens;CREATE TABLE movielens.movies (`  
`movieID IN PRIMARY KEY NOT NULL,`  
`title TEXT,releaseDate DATE`

`);`

### **Update the mysql.conf file**

`sudo cat /etc/logstash/conf.d/mysql.conf`

### **Mysql.conf**

```
input{
  jdbc{
    jdbc_connection_string => "jdbc:mysql://localhost:3306/movielens"
    jdbc_user => "student"
    jdbc_password => "*****"
    jdbc_driver_library => "home/student/usr/share/logstash/mysql-connector-java-8.0.16/mysql-connector-java-8.0.16.jar"
    jdbc_driver_class => "com.mysql.jdbc.Driver"
    statement => "SELECT * from movies"
  }
}

output{
  stdout { codec => json_lines }
  elasticsearch{
    hosts => ["localhost:9200"]
    index => "movielens-sql"
```



```
}
```

```
}
```

### **Run the commands:**

```
sudo /usr/share/logstash/bin/logstash -f /etc/logstash/conf.d/mysql.conf
```

```
curl -XGET 127.0.0.1:9200/movielens-sql/_search?pretty
```

### **Importing Data from .Csv:**

```
sudo cat /etc/logstash/conf.d/csv-read-drop.conf
```

#### **csv-read-drop.conf:**

```
input {
```

```
  file {
```

```
    path => "/home/student/csv-data/csv-schema-short-numerical.csv"
```

```
    start_position => "beginning"
```

```
  }
```

```
}
```

```
filter {
```

```
  csv {
```

```
    separator => ","
```

```
    skip_header => "true"
```

```
    columns =>
```

```
["id", "timestamp", "paymentType", "name", "gender", "ip_address", "purpose", "country", "age"]
```

```
  }
```

```
}
```

```
output {
```

```
  elasticsearch {
```

```
    hosts => "http://localhost:9200"
```

```
    index => "demo-csv"
```

```
  }
```

```
stdout {}
```

### **Run the Commands:**

```
sudo /usr/share/logstash/bin/logstash -f /etc/logstash/conf.d/demo-csv.conf
```

```
curl -XGET 127.0.0.1:9200/demo-csv/_search?pretty
```

### **Importing Data from .Json/.Log:**

- `cd /etc/logstash/conf.d/`
- `sudo vi json-read.conf`

### **json-read.conf: (using filter)**

```
input {  
  file {  
    start_position => "beginning"  
    path=> "/home/student/json-data/sample-json.log"  
  }  
}  
filter {  
  json {  
    source => "message"  
  }  
}  
output {  
  elasticsearch {  
    hosts => "http://localhost:9200"  
    index=>"demo-json"  
  }  
  stdout {}  
}
```

### **Run the commands:**

```
sudo /usr/share/logstash/bin/logstash -f /etc/logstash/conf.d/demo-json.conf
```

```
curl -XGET 127.0.0.1/demo-json/_search?pretty
```

**demo-json-drop.conf: (using mutate and removing unwanted fields)**

```
input {
  file {
    start_position => "beginning"
    path=> "/home/student/json-data/sample-json.log"
  }
}
filter {
  json {
    source => "message"
  }
  if [paymentType] == "Mastercard" {
    drop{}
  }
  mutate{
    remove_field =>["message","@timestamp","path","host","@version"]
  }
}
output {
  elasticsearch {
    hosts => "http://localhost:9200"
    index=>"demo-json-drop"
  }
  stdout {}
}
```

**Run the Commands:**

```
sudo /usr/share/logstash/bin/logstash -f /etc/logstash/conf.d/json-drop.conf
```

```
curl -XGET 'http://localhost:9200/demo-json-drop/_search?pretty'
```

## AGGREGATION

- Metrics - Avg,Min,Max
- Buckets - Histogram,Piechart

### Metrics:

### 1. Aggregation on ‘Ratings’ index

```
curl -XGET '127.0.0.1:9200/ratings/_search?pretty' -d '{
  "aggs" : {
    "ratings" : {
      "terms" : {
        "field": "rating"
      }
    }
  }
}'
```

```
"aggregations" : {
  "ratings" : {
    "doc_count_error_upper_bound" : 0,
    "sum_other_doc_count" : 0,
    "buckets" : [
      {
        "key" : 4.0,
        "doc_count" : 26818
      },
      {
        "key" : 3.0,
        "doc_count" : 20047
      },
      {
        "key" : 5.0,
        "doc_count" : 13211
      },
      {
        "key" : 3.5,
        "doc_count" : 13136
      },
      {
        "key" : 4.5,
        "doc_count" : 8551
      },
      {
        "key" : 2.0,
        "doc_count" : 7551
      },
      {
        "key" : 2.5,
        "doc_count" : 5550
      },
      {
        "key" : 1.0,
        "doc_count" : 2811
      },
    ]
  }
}
```

## 2. Using match and aggregation

```
curl -XGET '127.0.0.1:9200/ratings/_search?pretty' -d '
```

```
{
  "query":
    {"match":{"rating":5.0}},
  "aggs": {
    "ratings": {
      "terms": {
        "field":"rating"
      }
    }
  }
}
```

```
"aggregations" : {
  "ratings" : {
    "doc_count_error_upper_bound" : 0,
    "sum_other_doc_count" : 0,
    "buckets" : [
      {
        "key" : 5.0,
        "doc_count" : 13211
      }
    ]
  }
}
```

## 3. Using match\_phrase and avg\_aggregation

```
curl -XGET '127.0.0.1:9200/ratings/_search?pretty' -d '
```

```
{
  "query":{
    "match_phrase":{"title":"Star Wars"}},
  "aggs": {
    "avg_ratings": {
      "avg": {
        "field":"rating"
      }
    }
  }
}
```

```
"aggregations" : {
  "avg_ratings" : {
    "value" : 3.8587570621468927
  }
}
```

## **Buckets-Histogram:**

### **1. Histogram on the field ratings at interval 1.0**

```
curl -XGET '127.0.0.1:9200/ratings/_search?pretty' -d '{
  "aggs": {
    "whole_ratings": {
      "histogram": {
        "field": "rating", "interval": 1.0
      }
    }
  }
}
```

```
"aggregations" : {
  "whole_ratings" : {
    "buckets" : [
      {
        "key" : 0.0,
        "doc_count" : 1370
      },
      {
        "key" : 1.0,
        "doc_count" : 4602
      },
      {
        "key" : 2.0,
        "doc_count" : 13101
      },
      {
        "key" : 3.0,
        "doc_count" : 33183
      },
      {
        "key" : 4.0,
        "doc_count" : 35369
      },
      {
        "key" : 5.0,
        "doc_count" : 13211
      }
    ]
  }
}
```

## **Buckets-Time Series**

1. Histogram for time series data(calendar interval or fixed interval)

curl -XGET '127.0.0.1:9200/demo-grok/\_search?pretty' -d '

```
{
  "aggs": {
    "timestamp": {
      "date_histogram": {
        "field": "@timestamp", "fixed_interval": "5ms"
      }
    }
  }
}
```

```
"aggregations" : {
  "timestamp" : {
    "buckets" : [
      {
        "key_as_string" : "2024-07-22T07:02:13.990Z",
        "key" : 1721631733990,
        "doc_count" : 2
      },
      {
        "key_as_string" : "2024-07-22T07:02:13.995Z",
        "key" : 1721631733995,
        "doc_count" : 0
      },
      {
        "key_as_string" : "2024-07-22T07:02:14.000Z",
        "key" : 1721631734000,
        "doc_count" : 0
      },
      {
        "key_as_string" : "2024-07-22T07:02:14.005Z",
        "key" : 1721631734005,
        "doc_count" : 4
      },
      {
        "key_as_string" : "2024-07-22T07:02:14.010Z",
        "key" : 1721631734010,
        "doc_count" : 2
      }
    ]
  }
}
```

## **NESTED AGGREGATION:**

Average movie ratings that contain word 'Star' in its title

```
curl -XGET '127.0.0.1:9200/ratings/_search?pretty' -d '
```

```
{
  "query":
    {
      "match_phrase" : { "title" : "Star Wars" }},
      "aggs" :
        { "titles":
          {
            "terms" : { "field" : "title.raw" },
            "aggs" :
              { "avg_ratings" :
                { "avg" :
                  { "field" : "rating" }
                }
              }
          }
        }
    }
  }
}
```

## **KIBANA**

### **Install and enable Kibana:**

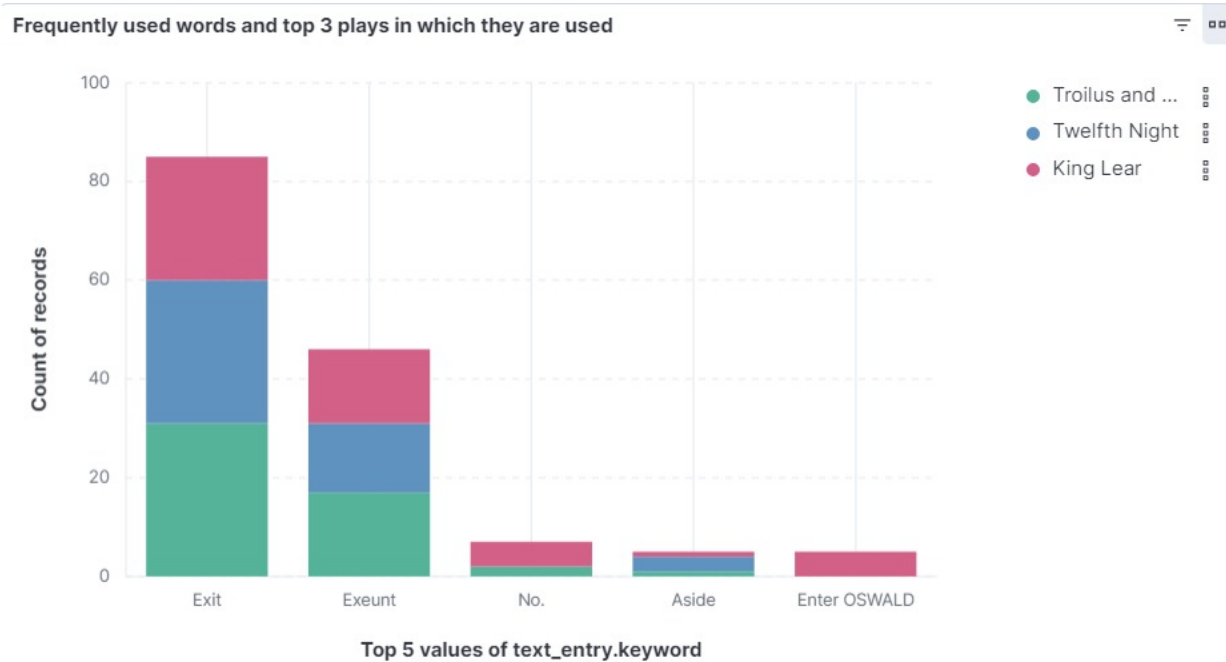
- sudo apt list kibana
- sudo apt-get install kibana=8.14.3
- sudo nano /etc/kibana/kibana.yml
- sudo /bin/systemctl enable kibana.service
- sudo /bin/systemctl start kibana.service
- Open kibana service from port localhost:5601



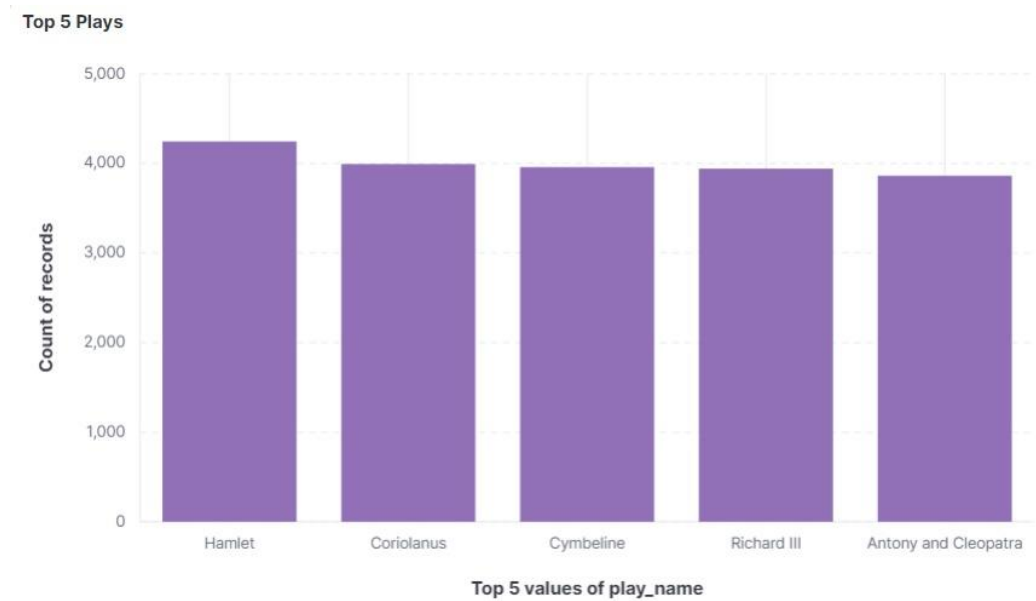
## Displaying the most frequently used words in the Shakespear's works using Tag Cloud in Aggregation



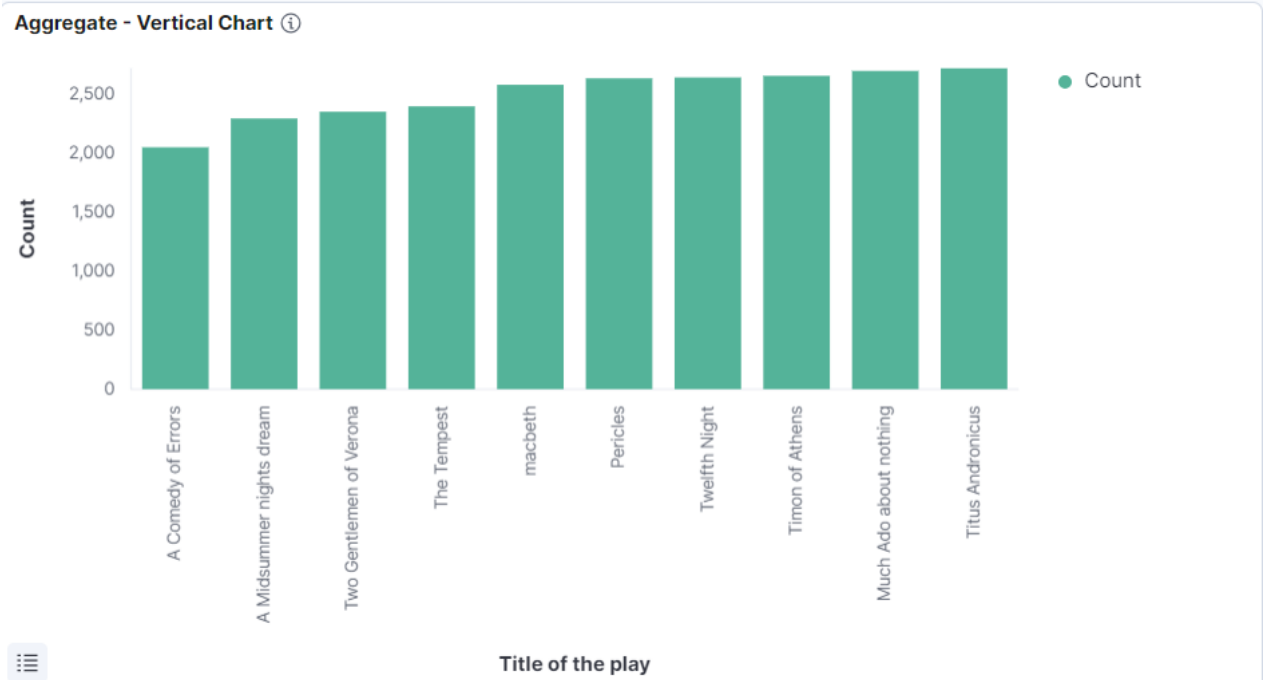
### Top 5 frequently used words and the plays in which they are used.



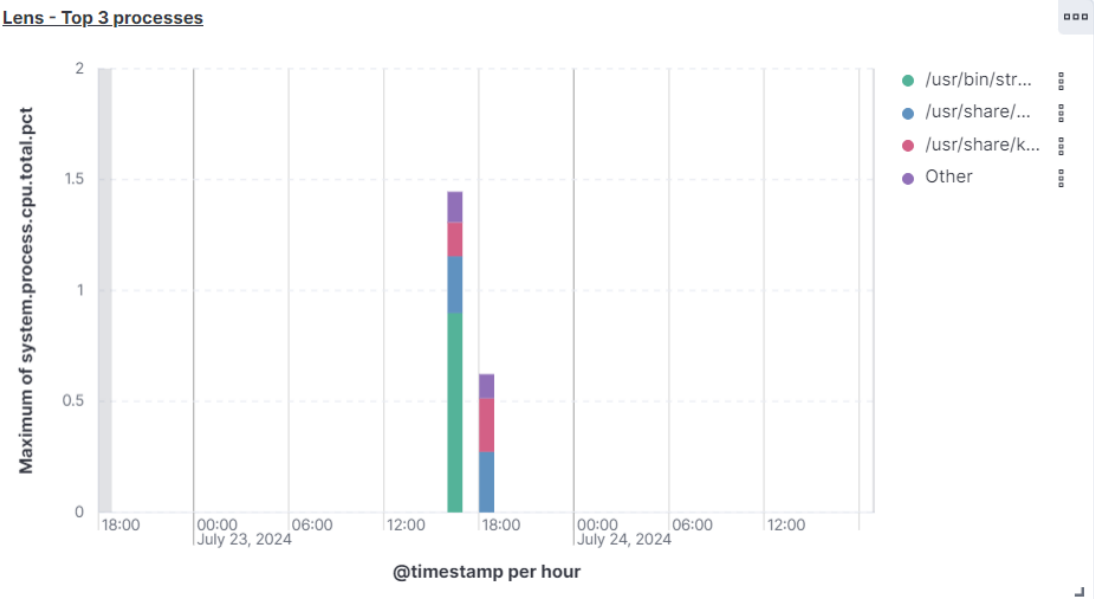
Vertical Bar chart to display the top 5 plays (having highest documents).



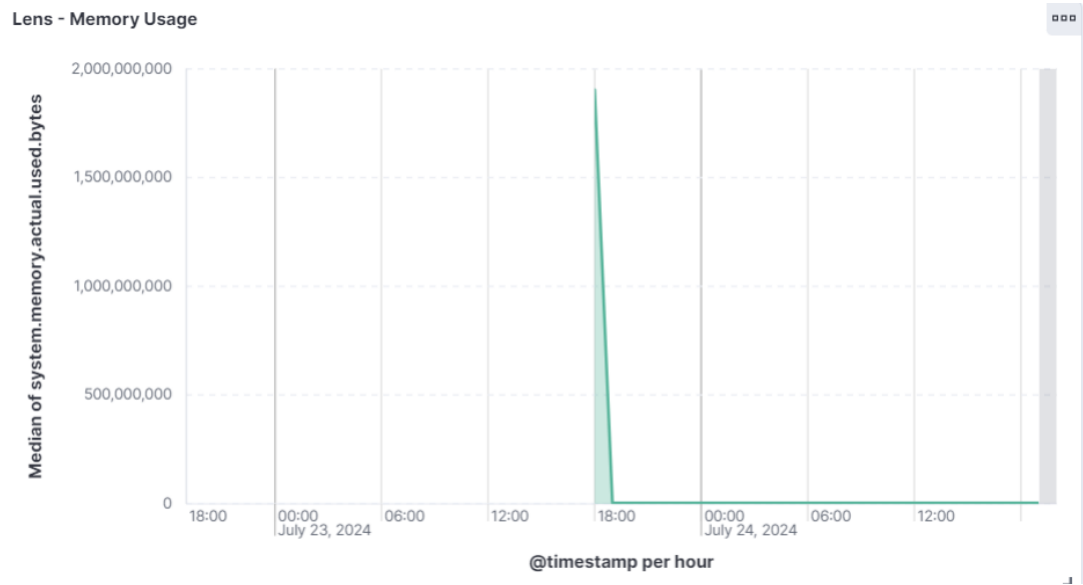
Vertical Bar chart displaying number of documents in last 10 plays.



# Top 3 processes

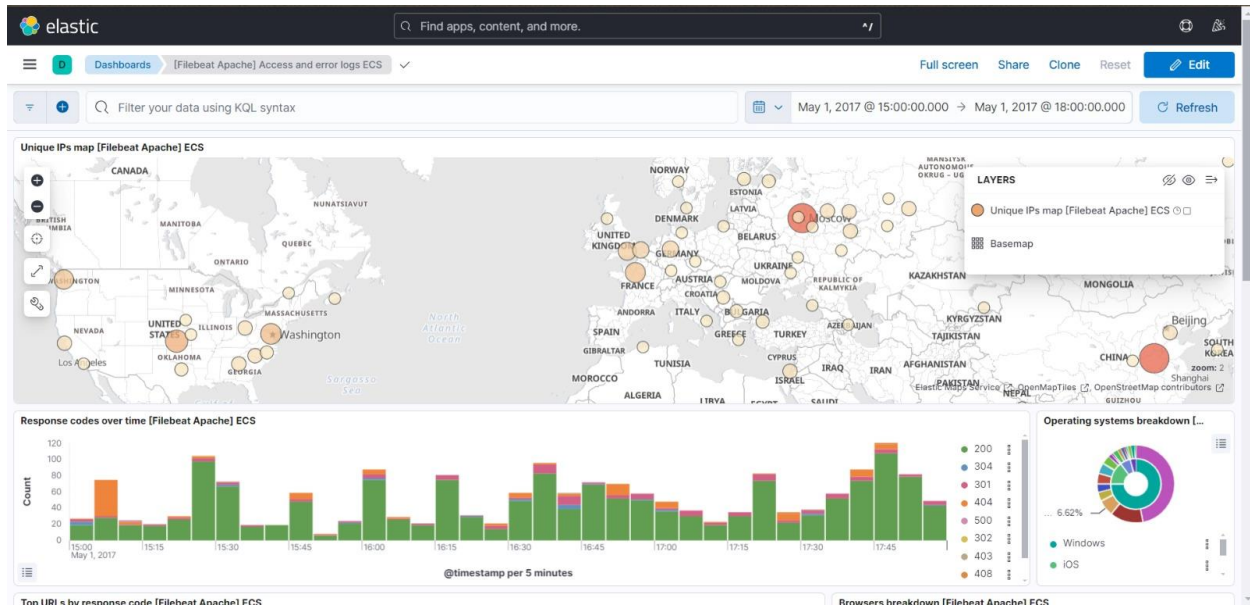


# Memory Usage



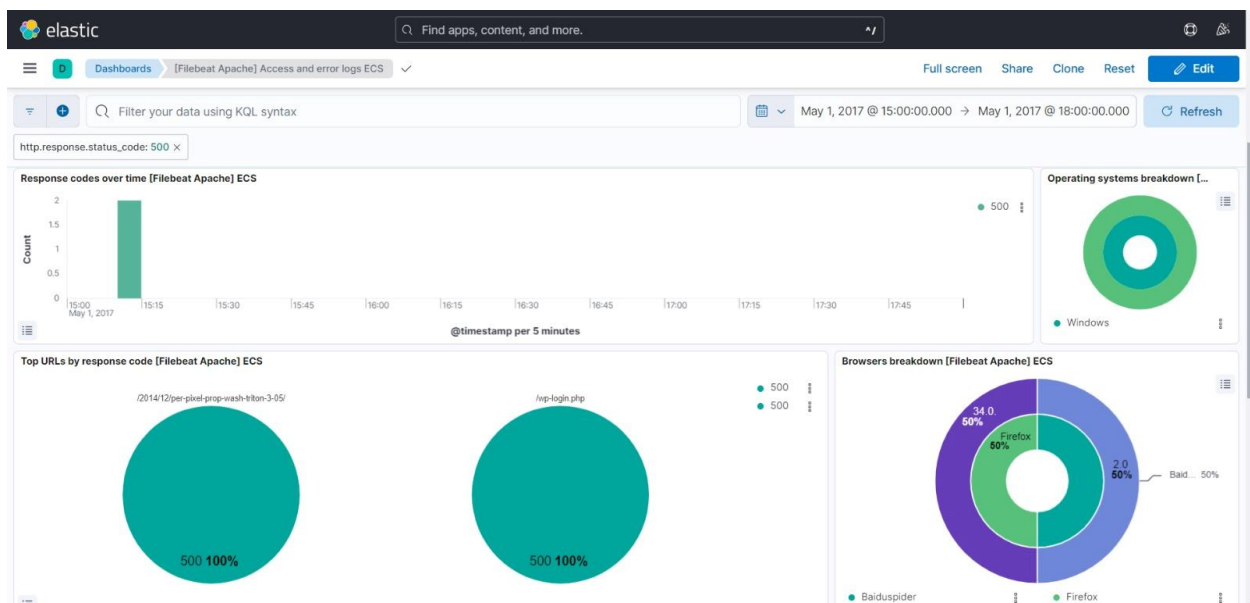
## LOG ANALYSIS IN KIBANA

The following dashboard displays the entire logs details of specific web



The following dashboard displays the logs details of specific web where it faced internal server error

**STATUS CODE: 500**





Documents (4)Field statistics

Columns 3Sort fields 1

Get the best look at your search results

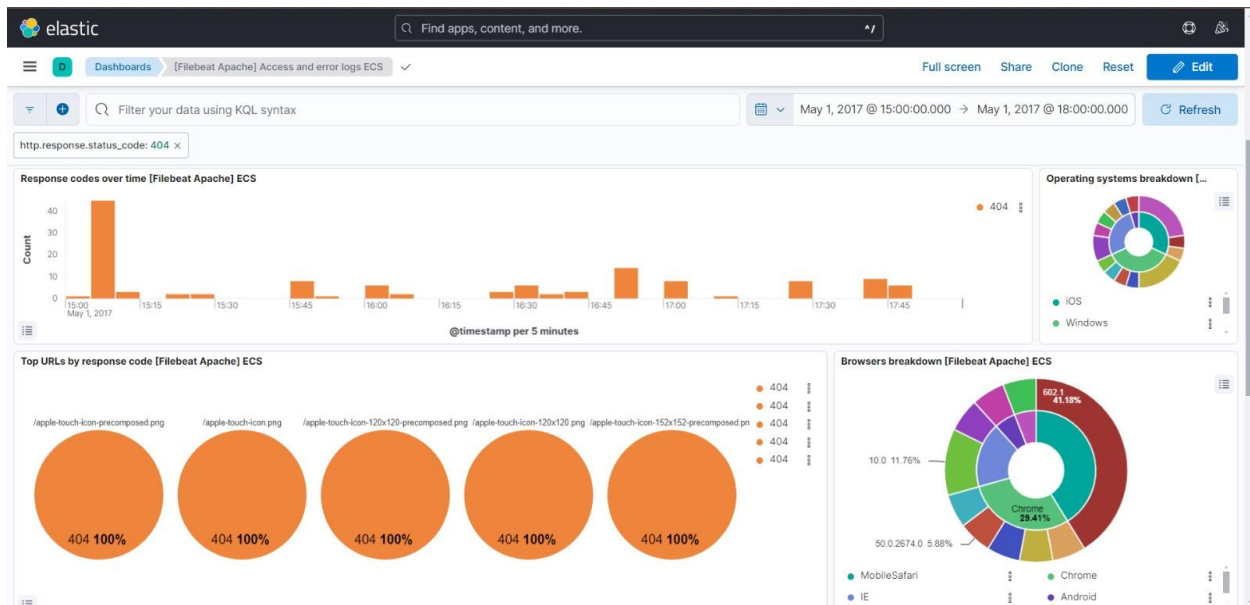
Add relevant fields, reorder and sort columns, resize rows, and more in the document table.

Take the tourDismiss

| @timestamp                                          | source.geo.location    | user_agent.name |
|-----------------------------------------------------|------------------------|-----------------|
| <input type="checkbox"/> May 4, 2017 @ 09:43:32.000 | POINT (-97.822 37.751) | Googlebot       |
| <input type="checkbox"/> May 4, 2017 @ 09:43:31.000 | POINT (-97.822 37.751) | Googlebot       |
| <input type="checkbox"/> May 4, 2017 @ 09:43:31.000 | POINT (-97.822 37.751) | Googlebot       |

The following dashboard displays the logs details of specific web where it faced resource not found error

**STATUS CODE: 404**



## ELASTICSEARCH and SQL

### 1. To get the type mappings

➤ `curl -XPOST 127.0.0.1:9200/_sql? Format=txt -d '`

`{"query": "DESCRIBE movies"}'`

| column        | type    | mapping |
|---------------|---------|---------|
| genre         | VARCHAR | text    |
| genre.keyword | VARCHAR | keyword |
| id            | VARCHAR | text    |
| id.keyword    | VARCHAR | keyword |
| title         | VARCHAR | text    |
| title.keyword | VARCHAR | keyword |
| year          | BIGINT  | long    |

### 2. To get the movies with year field less than 2000 and limit results to 10

➤ `curl -XPOST 127.0.0.1:9200/_sql?format=txt -d '`

```
{"query": "SELECT title, year from movies where year < 2000 limit 10" }
```

| title                       | year |
|-----------------------------|------|
| Toy Story                   | 1995 |
| Jumanji                     | 1995 |
| Grumpier Old Men            | 1995 |
| Waiting to Exhale           | 1995 |
| Father of the Bride Part II | 1995 |
| Heat                        | 1995 |
| Sabrina                     | 1995 |
| Tom and Huck                | 1995 |
| Sudden Death                | 1995 |
| GoldenEye                   | 1995 |

### CANVAS AND SQL

- Firstly, we must create a work pad which can consist of single or multiple pages.
- Each page can consist of elements like charts, graphs, maps, etc....

#### **Four elements:**

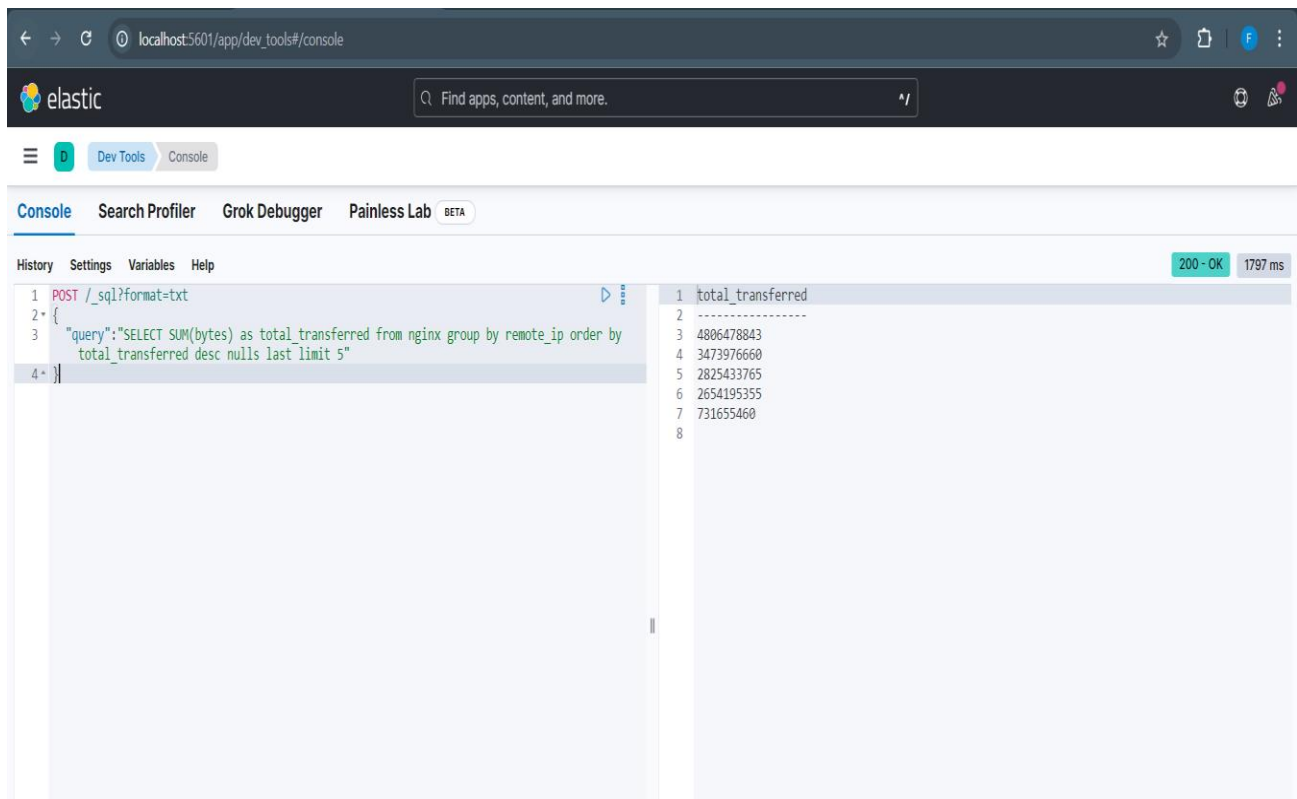
- Charts – bar chart, pie chart, doughnut, area, line, etc....
- Shapes – text boxes
- Images – can have no. of images varied based on the live data from elastic search
- Supporting elements – dropdown, filter options

#### **Canvas Data Sources**

Elasticsearch SQL queries

#### **Steps to create a canvas in Kibana**

- Inject Log data into Elastic search index - For examples: **nginx** in our canvas
- Select the Kibana space in which we want to work
- Click on Kibana dev tools and check if the SQL queries work fine on the index we will be using to create canvas metrics



1. Navigate to Analytics -> Canvas -> Work pad
2. After creating new work pad, start adding elements, for example, metric element inside chart



3. In display tab, we can change the font properties of the metric and in data source using Elasticsearch SQL
  - **SELECT Count (\*) AS count\_document FROM nginx.**
4. In display we will use the value **count\_document** to display the total logs in nginx index.
5. For inserting tables, insert the table element from charts and will get the data from elasticsearch SQL
  - **SELECT request, count (\*) as count\_requests FROM nginx GROUP BY request ORDER BY count\_requests DESC**
6. Insert chart -> Bar chart element. Change the data source (will be same as table)
 


In display, we will change the **x-axis to count\_requests** and **y-axis to request**.






## REQUEST STATS - NUMBER OF REQUESTS

| request #                          | count_requests # |
|------------------------------------|------------------|
| GET /downloads/product_1 HTTP/1.1  | 30272            |
| GET /downloads/product_2 HTTP/1.1  | 21034            |
| GET /downloads/product_3 HTTP/1.1  | 73               |
| HEAD /downloads/product_2 HTTP/1.1 | 70               |
| HEAD /downloads/product_1 HTTP/1.1 | 13               |



< 1 >

## TOP 5 IP addresses - TRANSFERRED BYTES



| remote_ip #    | bytes_transferred # |
|----------------|---------------------|
| 107.23.7.76    | 4806478843          |
| 114.80.245.62  | 3473976660          |
| 74.205.117.244 | 2825433765          |
| 54.239.240.49  | 2654195355          |
| 54.208.16.21   | 731655460           |

107.23.7.76 114.80.245.62 74.205.117.244 54.239.240.49 54.208.16.21

< 1 >

51,462  
Logs

33,939,67  
Bytes Transferred

2,659  
Unique Ip's

136  
Agents



## BACKUP AND TROUBLE SHOOTING

### Categories

- Node setup
- Discovery and cluster formation
- Indexing data and sharding
- Searching
- Backing up data

### Steps to perform Back up in Elasticsearch:

- `sudo nano /etc/elasticsearch/elasticsearch.yml`
- Now add `path.repo : ["/home/student/backups"]` after the `path.logs` in `elasticsearch.yml`
- `sudo cp /etc/elasticsearch/elasticsearch.yml ~/`
- `sudo mkdir -p /home/student/backups`
- `sudo chgrp elasticsearch /home/student/backups`
- `sudo chmod g+w /home/student/backups/`

- `sudo /bin/systemctl stop elasticsearch.service`
- `sudo /bin/systemctl start elasticsearch.service`
- `curl --request PUT localhost:9200/_snapshot/backup-repo \`  
`--data-raw ' {`  
`"type" : "fs",`  
`"settings": {`  
`"location":"/home/student/backups/backup-repo"`  
`}`  
`}'`
- `curl --request PUT localhost:9200/_snapshot/backup-repo/snapshot-1`
- `curl --request GET localhost:9200/_snapshot/backup-repo/snapshot-1? Pretty`

### **Potential Issues and Trouble Shooting in Elasticsearch:**

#### **Open a new PUTTY window – Terminal 2**

`sudo visudo`

In the bottom of the file add:username ALL=(elasticsearch) NOPASSWD: ALL

`sudo -su elasticsearch`

`cd /var/log/elasticsearch/`

`tail -n 500 elasticsearch.log | grep ERROR`

`cat Elasticsearch.log | grep Bootstrap --context=3`

### **MEMORY LOCK ISSUE:**

#### **In Terminal-2:**

- `sudo nano /etc/elasticsearch/elasticsearch.yml`
- Uncomment the line `bootstrap.memory_lock:true`
- `sudo systemctl stop elasticsearch.service`
- `sudo systemctl start elasticsearch.service`

**\*\*It will throw you an error: It will show memory is locked error**

### Go back to Terminal-1

- `sudo systemctl edit elasticsearch.service`
- Add the following in the file to resolve the above error:  
[Service]  
LimitMEMLOCK=infinity
- `sudo systemctl start elasticsearch.service`

### HEAP MEMORY ALLOCATION ISSUE:

#### In Terminal-2:

- `sudo nano /etc/elasticsearch/jvm.options`
- Comment out both `-Xmslg` and `Xmxlg` in the `jvm` file
- And add:  
    `-Xms500m`  
    `-Xmslg`
- `sudo systemctl stop elasticsearch.service`
- `sudo systemctl start elasticsearch.service`

**\*\*It will throw error initial heap size not equal to the initial allocation error**

- `sudo nano /etc/elasticsearch/jvm.options`
- Uncomment out both `-Xmslg` and `Xmxlg`
- And remove:  
    `-Xms500m`  
    `-Xmslg`
- `sudo systemctl stop elasticsearch.service`
- `sudo systemctl start elasticsearch.service`

### NODE SETUP ISSUES:

#### In Terminal-2

- `sudo cat /usr/lib/systemd/system/elasticsearch.service`
- `sudo nano /etc/elasticsearch/elasticsearch.yml`
- Comment out:  
    `discovery.seed_hosts: ["127.0.0.1"]` and  
    `cluster.initial_master_nodes: ["node-1"]`

**\*\*It will throw master not found exception**

- `sudo systemctl stop elasticsearch.service`

### In Terminal-1

- `rm -rf /var/lib/elasticsearch/*`
- Go back to terminal 2 and do
- `sudo vim /etc/elasticsearch/elasticsearch.yml`
- Uncomment and change:  
    `cluster-name:lecture-cluster`  
    `discovery.seed_hosts: ["127.0.0.1:9301"]`  
    `sudo systemctl start elasticsearch.service`

**\*\*The cluster\_uuid will be na which means the cluster is not formed.**

`sudo systemctl stop elasticsearch.service`

### Reasons for not forming cluster:

**It may be due to network issues where nodes within cluster might be unable to communicated with each other.**

### Creating index with 1 shard and 1 replica:

- `Curl -request PUT localhost:9200/test \`  
    `--data-raw '{`  
    `"settings":{`  
    `"number_of_shards":1,`  
    `"number_of_replicas":1}}`
- To check about the shard's status using:  
    1. `Curl localhost:9200/_cat/shards? V`  
    2. It will return that the status as started or unassigned.
- Cluster allocation API to explain why shards aren't allocated  
    1. `Curl localhost:9200/_cluster/allocation/explain? Pretty`  
    2. Reason – replica to the same node is not allowed.  
    3. How to overcome? Add a new node and take replica to the new node.

### Steps to setup 2<sup>nd</sup> node:

- `Sudo nano /etc/elasticsearch-node2/elasticsearch.yml`  
    `Node.name : node-2`  
    `Master.nodes will be node-1 and node-2`

- Start the 2<sup>nd</sup> node on the same VM  
Sudo systemctl start elasticsearch-node2
- To simulate red status
  1. Create index with 2 shards and no. of replicas to 0 (no backups)
  2. Curl –request PUT localhost:9200/test \
 

```
--data-raw '{
  "settings": {
    "number_of_shards":2,
    "number_of_replicas":0
  }
}'
```
  3. Two shards will be in 2<sup>nd</sup> nodes in the cluster
- We will stop node-2 and cut communication between 2 nodes in the cluster
- Now if we check status of cluster, it will be red  
Curl –silent localhost:9200/\_cluster/health?pretty | grep status
- To overcome this, we can allocate a new shard  
Curl –request POST “localhost:9200/\_cluster/reroute?pretty” \
 

```
--data-raw '{
  "commands": {
    {
      "allocate_empty_primary": {
        "index": "test",
        "shard": 1,
        "node": "node-1"
      },
      "accept_data_loss": "true"
    }
  }
}'
```
- Since our data is lost, we will be using backup data
- Move configuration file to previous place (default setting)

## **INDEX DESIGN CHANGES (SPLITTING, SHRINKING)**

### **Index settings:**

- ➔ Dynamic – can be changed after index creation
  - Number\_of\_replicas
  - Refresh intervals
  - Blocks – disabling readability/writability of index
  - Pipeline – preprocessing pipeline for every documents

➔ Static – can't be changed after index creation

Number\_of\_shards

### **Sharding goals:**

High availability - working uninterrupted for a long time

High resiliency – resist errors

### **Increase/decrease shards**

#### **To decrease**

➤ POST /{source\_index}/\_shrink/{target\_index-name}

#### **To increase**

➤ POST /{source\_index}/\_split/{target\_index-name}

## **SNAPSHOTS**

- Backups to NAS, Amazon S3, Azure
- Store only changes
- Elasticsearch.yml
- Path.repo : ["/home/<user>/backups"]

### **Setup Steps:**

1. Configuration file
  - Path.repo : ["/home/student/backups"]
2. Mkdir backups
3. Restart elasticsearch.service
4. PUT \_snapshot/backup-repo

```
{  
  "type": "fs",  
  "settings": {  
    "location": "/home/student/backups/backup-repo"  
  }  
}
```

## SPRINGBOOT AND ELASTIC STACK

### Spring boot-Elastic Stack

Spring Boot application sending logs to "Elastic Search Logstash Kibana" demo

### Why Elastic Stack?

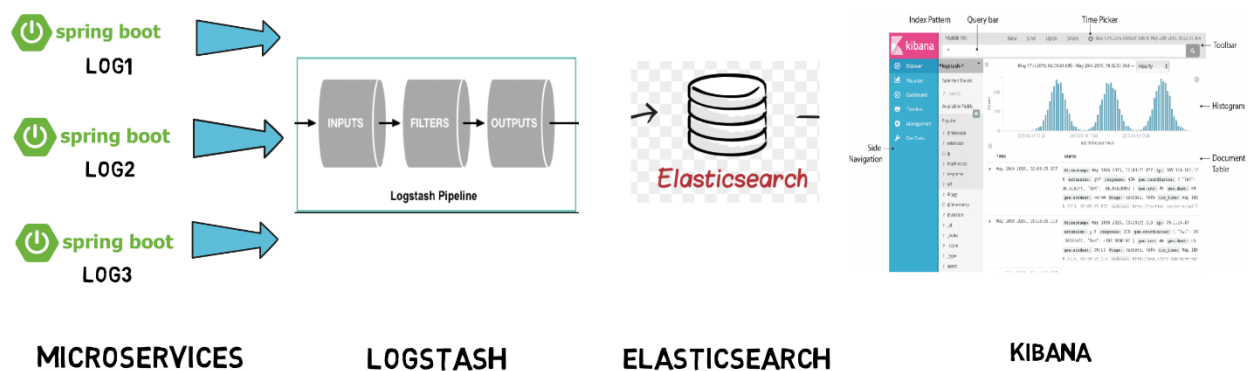
With use of microservices, we have been able to overcome many traditional problems and it allow us to create stable distributed applications with desired control on the code, team size, maintenance, release cycle, cloud ennoblement, automation, etc. But it has also introduced few challenges in other areas e.g. "distributed log management" and ability to view logs of full transaction distributed among many services and distributed debugging in general.

Actually, the challenge is that microservices are isolated among themselves and they do not share common database and log files. As the number of microservice increases and we enable cloud deployment with automated continuous integration tools, it is very much necessary to have some provision of debugging the components when we have any problem.

Thanks to the open source. We already have bundle of tools which can do the magic if used properly together. One such popular set of tools are Elastic Search, Logstash and Kibana – together referred as ELK stack. They are used for searching, analyzing, and visualizing log data in a real time.

### ELK Stack with SPRINGBOOT & MICROSERVICES:

## ELK – ELASTIC SEARCH – LOGSTASH – KIBANA – ARCHITECTURE



Logstash processes the application log files based on the filter criteria we set and sends those logs to Elasticsearch. Through Kibana, we view and analyze those logs when required. Elasticsearch is a distributed, JSON-based search and analytics engine designed for horizontal scalability, maximum reliability, and easy management.

Logstash is a dynamic data collection pipeline with an extensible plugin ecosystem and strong Elasticsearch synergy.

Kibana gives the visualization of data through a UI.

### **Elastic stack configuration**

All these three tools are based on JVM and before start installing them, please verify that JDK has been properly configured. Check that standard JDK 1.8 installation, JAVA\_HOME and PATH set up is already done.

### **Elasticsearch**

Download latest version of Elasticsearch from this download page <https://www.elastic.co/downloads/elasticsearch> and unzip it any folder.

Run bin\elasticsearch.bat from command prompt.

By default, it would start at <http://localhost:9200>

### **Kibana**

Download the latest distribution from download page

<https://www.elastic.co/de/downloads/kibana> and unzip into any folder.

Open config/kibana.yml in an editor and set elasticsearch.url to point at your Elasticsearch instance. In our case as we will use the local instance just uncomment elasticsearch.hosts:

"<http://localhost:9200>"

Run bin\kibana.bat from command prompt.

Once started successfully, Kibana will start on default port 5601 and Kibana UI will be available at <http://localhost:5601>

### **Logstash**

Download the latest distribution from download page <https://www.elastic.co/downloads/logstash> and unzip into any folder.

Create one file logstash.conf as per configuration instructions.

Now run bin/logstash -f logstash.conf to start logstash

### **Logstash configuration**

We need to create a logstash configuration file so that it listen to the log file and push log messages to elastic search.

---

```
input {
  file {
    type => "java"
    path => "<PLEASE_UPDATE_YOUR_LOG_FILE_PATH_HERE>"
    codec => multiline {
      pattern => "^%{YEAR}-%{MONTHNUM}-%{MONTHDAY} %{TIME}.*"
      negate => "true"
      what => "previous"
    }
  }
}
```



```

    }
  }

filter {
  #If log line contains tab character followed by 'at' then we will tag that entry as stacktrace
  if [message] =~ "\tat" {
    grok {
      match => ["message", "^(\\tat)"]
      add_tag => ["stacktrace"]
    }
  }

  grok {
    match => [ "message",
      "(?<timestamp>%{YEAR}-%{MONTHNUM}-%{MONTHDAY}
%{TIME}) %{LOGLEVEL:level} %{NUMBER:pid} --- \\[(?<thread>[A-Za-z0-9-]+)\\] [A-Za-
z0-9.]*\\.?(?<class>[A-Za-z0-9#_]+)\\s*:\\s+(?<logmessage>.*)",
      "message",
      "(?<timestamp>%{YEAR}-%{MONTHNUM}-%{MONTHDAY}
%{TIME}) %{LOGLEVEL:level} %{NUMBER:pid} --- .+? :\\s+(?<logmessage>.*)"
    ]
  }

  date {
    match => [ "timestamp" , "yyyy-MM-dd HH:mm:ss.SSS" ]
  }
}

output {

  stdout {
    codec => rubydebug
  }

  # Sending properly parsed log events to elasticsearch
  elasticsearch {
    hosts => ["localhost:9200"]
  }
}
...

```

Once ELK stack is up and running. You could start this demo project and call the REST endpoints, view the logs in KIBANA.

Do a final maven build using `mvn clean install` and start the application and test by browsing <http://localhost:8080/elk> or

<http://localhost:8080/elkdemo> or <http://localhost:8080/exception>

Don't be afraid by seeing the big stack trace in the screen as it has been done intentionally to see how ELK handles exception message.

Go to the application root directory and verify that the log file i.e. elk-example.log has been created and do a couple of visits to the endpoints and verify that logs are getting added in the log file.

## Test The Demo Application with ELK

Now when all components are up and running, let's verify the whole ecosystem.

Go to application and test the end points couple of times so that logs got generated and then go to Kibana console and see that logs are properly stacked in the Kibana with lots of extra features like we can filter, see different graphs etc. in built.

Here is the view of generated logs in Kibana.

The screenshot shows the Kibana interface with the 'Discover' tab selected. The left sidebar lists various navigation options. The main content area displays a list of log entries. The first entry is a configuration document. The second entry is a log message with a detailed stack trace. The stack trace indicates a ServletException caused by a Request processing failed, with a nested exception of java.lang.IllegalArgumentException. The message field is expanded, showing the full stack trace details.

```
space.name: Default space.description: This is your default space! space.color: #00bfb3 space.reserved: true type: space updated_at: March 8th 2019, 09:56:36.752 _id: space:default _type: doc _index: .kibana_1 _score: 1

config.buildNum: 10513 config.defaultIndex: 161e84c0-4187-11e9-9fef-914edcb41c95 type: config updated_at: March 8th 2019, 10:46:46.715 _id: config:6.6.1 _type: doc _index: .kibana_1 _score: 1

index-pattern.title: * index-pattern.fields: [{"name":"_id","type":"string","count":0,"scripted":false,"searchable":true,"aggregatable":true,"readFromDocValues":false}, {"name":"_index","type":"string","count":0,"scripted":false,"searchable":true,"aggregatable":true,"readFromDocValues":false}, {"name":"_score","type":"number","count":0,"scripted":false,"searchable":false,"aggregatable":false,"readFromDocValues":false}, {"name":"_source","type":"_source","count":0,"scripted":false,"searchable":false,"aggregatable":false,"readFromDocValues":false}, {"name":"_type","type":"string","count":0,"scripted":false,"searchable":true,"aggregatable":true,"readFromDocValues":false}, {"name":"apm-telemetry.has_any_services","type":"boolean","count":0,"scripted":false,"searchable":true,"aggregatable":true,"readFromDocValues":false}]

host: Vinoths-MBP.fritz.box path: /Users/vinothkumar/workspace/ELK/elk-example.log @timestamp: 2019-03-08T09:49:14.811Z tags: multiline, stacktrace, _grokparsefailure @version: 1 type: java message: 2019-03-08 10:47:07.831 ERROR 11333 --- [http-nio-8080-exec-1] o.a.c.c.C.[.[/].[dispatcherServlet] : Servlet.service() for servlet [dispatcherServlet] in context with path [] threw exception [Request processing failed; nested exception is java.lang.IllegalArgumentException: Type must be a parameterized type: java.lang.Class] with root cause java.lang.IllegalArgumentException: Type must be a parameterized type: java.lang.Class at org.springframework.util.Assert.instanceCheckFailed(Assert.java:655) ~[spring-core-5.1.5.RELEASE.jar:5.1.5.RELEASE] at org.springframework.util.Assert.isInstanceOf(Assert.java:555) ~[spring-core-5.1.5.RELEASE.jar:5.1.5.RELEASE] at org.springframework.util.Assert.isInstanceOf(Assert.java:555) ~[spring-core-5.1.5.RELEASE.jar:5.1.5.RELEASE] at org.springframework.util.Assert.isInstanceOf(Assert.java:555) ~[spring-core-5.1.5.RELEASE.jar:5.1.5.RELEASE]
```

| @timestamp               | @version | _id                 | _index              | _score | _type | host                  | message                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|--------------------------|----------|---------------------|---------------------|--------|-------|-----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 2019-03-08T09:49:14.811Z | 1        | PRG0XGk8HoJmQ937Y41 | logstash-2019.03.08 | 1      | doc   | Vinoths-MBP.fritz.box | 2019-03-08 10:47:07.831 ERROR 11333 --- [http-nio-8080-exec-1] o.a.c.c.C.[.[/].[dispatcherServlet] : Servlet.service() for servlet [dispatcherServlet] in context with path [] threw exception [Request processing failed; nested exception is java.lang.IllegalArgumentException: Type must be a parameterized type: java.lang.Class] with root cause java.lang.IllegalArgumentException: Type must be a parameterized type: java.lang.Class at org.springframework.util.Assert.instanceCheckFailed(Assert.java:655) ~[spring-core-5.1.5.RELEASE.jar:5.1.5.RELEASE] at org.springframework.util.Assert.isInstanceOf(Assert.java:555) ~[spring-core-5.1.5.RELEASE.jar:5.1.5.RELEASE] at org.springframework.util.Assert.isInstanceOf(Assert.java:555) ~[spring-core-5.1.5.RELEASE.jar:5.1.5.RELEASE] at org.springframework.util.Assert.isInstanceOf(Assert.java:555) ~[spring-core-5.1.5.RELEASE.jar:5.1.5.RELEASE] |